

# Gioia Zheng

Rome, Italy — [gioia.zheng.stud@gmail.com](mailto:gioia.zheng.stud@gmail.com) — [linkedin.com/in/gioia-zheng-9233a0303](https://linkedin.com/in/gioia-zheng-9233a0303)  
[github.com/GioiaZheng](https://github.com/GioiaZheng)

## Profile

Undergraduate at Sapienza University of Rome. Building MS MARCO retrieval-augmented QA end-to-end on a single 6-core CPU machine — BM25, dense (SBERT/FAISS), cross-encoder rerank, T5 generation — with paired-bootstrap CIs, NLI-entailment grounding, and a regression failure taxonomy. Interested in the gap between surface-form and semantic-faithfulness metrics in RAG evaluation.

## Education

### Sapienza University of Rome

Sep 2022 – Present

B.Sc. in Applied Computer Science and Artificial Intelligence

## Experience

### NLP Algorithm Engineer Intern

Feb 2026 – May 2026

Microsoft — Natural Language Computing (Remote)

- Headline result on MS MARCO dev/small (6,980 queries, 8.8M-passage corpus): replacing BM25 first-stage with dense top-100 reranked by a cross-encoder, same frozen T5-small generator, lifted Token-F1 **0.197** → **0.368** and ROUGE-L **0.193** → **0.368**. 95% paired-bootstrap CIs on per-query  $\Delta$  (10k resamples, seed 42) strictly above 0 across all four surface-form metrics.
- Pipeline I built and own end-to-end: **bm25s** sparse (MRR@10 0.170, in line with Anserini's 0.184; gap consistent with tokenizer differences), dense **all-MiniLM-L6-v2** over a FAISS **IndexFlatIP** on L2-normalised embeddings, **cross-encoder/ms-marco-MiniLM-L-6-v2** over the dense top-100 (538k pairs at ~32 pairs/s, 4h37m wall-clock on 6 cores), T5-small generation.
- Ruled out a surface-form artefact with a DistilBERT-BERTScore proxy on 3,000 paired qids: semantic-proxy  $\Delta$  +0.173 sits within 0.002 of the surface-form Token-F1  $\Delta$ , well inside encoder noise. Chose DistilBERT over canonical **roberta-large** as an explicit CPU-cost vs. signal tradeoff, documented in the report.
- 40-query deterministic-rule triage of the 233-query regression bucket: **~90% of regressions are generator-side truncation** (mid-word or  $\leq 3$ -token outputs), not retrieval failure. Falsified the decode-budget reading by re-running full dev/small at **max\_new\_tokens=64→128**: truncation share moved 90% → 87.5%, all four rerank  $\Delta$ s moved  $< 0.005$ . The output style is intrinsic to T5-small on this prompt format, not a budget symptom.
- NLI-entailment grounding (**cross-encoder/nli-deberta-v3-small**, 3,000-pair subsample) is the only metric whose paired  $\Delta$  reverses sign:  $-0.145$  (95% CI  $[-0.16, -0.13]$ ,  $p < 0.001$ ). Hypothesis: fragmentary mid-word outputs that inflate token overlap cannot be entailed by a sentence-level NLI cross-encoder. Flagged as a generator-swap follow-up rather than a regression to roll back.
- Reproducibility scaffold: per-run manifest (git SHA, config hash, dependency-file hashes, environment fingerprint), pinned **requirements-lock.txt**, ~295-test **pytest** suite, CI on every push, batch-eval design with no hidden one-shot CLI.

## Projects

### Handwritten OCR pipeline

[github.com/GioiaZheng/handwritten-ocr-system](https://github.com/GioiaZheng/handwritten-ocr-system)

PyTorch — residual CNN + BiLSTM + CTC, IAM dataset

- CRNN for variable-length handwritten line transcription: 9 residual conv blocks (32→128 channels, batchnorm + leaky-ReLU + dropout 0.2), 2 stacked BiLSTM layers (hidden 256 / 64), linear head over vocab + CTC blank.
- Training and held-out evaluation with per-batch CER and WER (**torchmetrics**) streamed to TensorBoard; greedy CTC decode at inference; preprocessing pipeline (grayscale, normalisation, augmentation).

### LeetCode solutions and pattern notes

[github.com/GioiaZheng/Leetcode-Solutions](https://github.com/GioiaZheng/Leetcode-Solutions)

Python — schema-validated repo, GitHub Actions CI

- 98 problems, one directory each, with per-problem metadata (id, difficulty, topics, status) validated against a JSON schema. Topic index and catalog regenerated on every push; CI fails if the generated docs drift out of sync with the source metadata.
- Quality gate on every push: **ruff** lint, **pytest** on a curated subset, structural validation, matrix CI on Python 3.11 and 3.12.

### Chat service

[github.com/GioiaZheng/go-chat-system](https://github.com/GioiaZheng/go-chat-system)

Go backend + Vue frontend + SQLite

- Go REST backend with SQLite for single-writer durability; Vue frontend embedded into the Go binary at build time so the service ships as one static artifact.
- OpenAPI schema checked in CI; per-push matrix: **gofmt**, **go vet**, unit tests, OpenAPI lint, frontend build.

## Skills

### Primary

Python, PyTorch, **bm25s**, FAISS, sentence-transformers, cross-encoder reranking, paired-bootstrap CIs, BERTScore / NLI-entailment grounding, **pytest**

### Working

Go, Docker, Linux, GitHub Actions CI, OpenAPI, SQLite, LaTeX

### Familiar

JavaScript / React Native, Vue

### Languages

English (Professional), Italian (Bilingual), Chinese (Native), German (Beginner)